

ISHA AI

Intelligent Speech & Haptic Assistant

"Hey ISHA" — Two words. One billion people. Zero cloud. Zero cost. Forever.

Complete Technical Documentation

Version 1.0 | Engineering Master Document | Confidential
May 2026 | ISHA AI Engineering Team

Metric	Value
Document Sections	31
Architecture Depth	Staff Engineer / CTO Level
Primary Platform	Android (Lite/Standard) + iOS (Pro)
Target Users	1 Billion+ Indian & Global Smartphone Users
Core Philosophy	Offline-First. Privacy-Absolute. Free Forever.
Wake Word	Hey ISHA

Table of Contents

Table of Contents.....2

1. Executive Summary4

 What Is ISHA AI?4

 Who It Helps.....4

 The Core Insight.....4

 Five Laws of ISHA.....4

2. Problem Statement — Deep Analysis5

 Problem 1: Cloud Dependency.....5

 Problem 2: Language Exclusion5

 Problem 3: Hallucination in High-Stakes Contexts.....5

 Problem 4: Stateless AI — No Memory5

 Problem 5: Hardware Inefficiency5

 Problem 6: The Cost Barrier6

3. Existing Solutions Comparison.....7

 Feature Comparison Matrix7

 Why Each Competitor Fails ISHA's Market.....7

 ChatGPT.....7

 Google Gemini.....7

 Apple Intelligence8

4. Full System Architecture9

 High-Level Architecture9

 Complete Request Lifecycle — 20 Steps.....9

5. Model Architecture & AI Stack.....11

 The Four-Tier Model System Overview11

 Why Qwen2.5-1.5B for Tier 1 — The Critical Decision11

 Inference Engines11

 Embedding & Retrieval Models12

6. Voice & Wake Word System13

 Two-Stage Wake Detection Architecture13

 Speech-to-Text — Whisper.cpp.....13

 Text-to-Speech — Kokoro-82M13

7. Vision & Live Camera Architecture.....15

 Sparse Event-Driven Vision — The Core Design.....15

 Vision Capabilities by Tier15

8. Knowledge Pack System17

 Official Pack Catalog17

 Pack Architecture — Internal Structure.....17

 Install-Time Pre-Processing.....18

 Community & Enterprise Packs18

9. Memory Architecture	19
Four-Layer Memory Hierarchy	19
Semantic Compression — The Core Innovation	19
24-Hour Chat Lifecycle	19
10. Retrieval Engine (RAG)	21
The Retrieval Pipeline	21
Online Mode — MCP + Live RAG	21
11. ISHA Mesh — Peer Intelligence Network	22
The Concept	22
Mesh Protocol Stack	22
Privacy in the Mesh — Critical Design	22
12. Security & Privacy Deep Dive	23
Threat Model	23
What ISHA Cannot Do — By Architecture	23
13. Performance Engineering	24
Latency Budget	24
Thermal State Machine	24
14. End-to-End User Journey	25
Ramesh — Wheat Farmer, Vidarbha, 7:15 AM	25
15. Technical Tradeoff Analysis	26
16. Database Architecture	27
Storage Layer Overview	27
Key SQLite Tables	27
17. Platform Architecture — iOS vs Android	28
18. Hardware Tiering System	29
Device Capability Detection	29
19. Infrastructure & Deployment	30
Server Infrastructure — Minimal by Design	30
Cost Comparison	30
20. Testing Strategy	31
21. Engineering Mentor Review	32
What Is Excellent	32
What Needs Work — Honest Assessment	32
Hidden Risks	33
22. Production Readiness Assessment	34
23. Product Roadmap	35
24. Failure Scenarios & Recovery	36
25. Glossary	37

1. Executive Summary

What Is ISHA AI?

ISHA AI (Intelligent Speech & Haptic Assistant) is a fully offline-capable, privacy-first AI cognitive runtime designed to run natively on smartphones — from a ₹15,000 Android device to a flagship iPhone — without requiring internet connectivity, cloud subscriptions, or any ongoing payment.

ISHA is not a chatbot. It is not an app that calls an API. It is an on-device cognitive operating layer — a persistent intelligence system that lives inside your phone, understands your context, speaks your language, sees through your camera, and thinks using your hardware. Entirely.

Who It Helps

User	Problem ISHA Solves
Student in Bihar	Offline JEE/NEET prep in Hindi — no coaching center needed
Farmer in Vidarbha	Crop disease identification via camera, fully offline
Woman in rural Rajasthan	Legal rights guidance, private — no one sees her query
Elderly person in Kerala	Voice-first Malayalam interface, no typing required
Developer in Bangalore	Privacy-preserving coding assistant — no data leaks
Family with 3 phones	Shared ISHA Mesh intelligence — better than cloud AI

The Core Insight

Every other AI: Cloud → Model → Answer | ISHA: Device → Structure → Retrieval → Small Model → Answer

A well-organized 1.5B parameter model with perfect retrieval beats a 70B cloud model with no context. ISHA proves this.

Five Laws of ISHA

Law	Principle	Meaning
Law 1	Offline First	Every feature works without internet. Online is an enhancement, never a requirement.
Law 2	Privacy Absolute	No user data ever leaves the device by default — not to ISHA servers, not to anyone.
Law 3	Hardware = Subscription	Your device capability is your entitlement. No payment gates intelligence.
Law 4	Language Native	Hindi, Tamil, Telugu, Bengali, Marathi are first-class citizens — not afterthoughts.
Law 5	Honest Intelligence	ISHA never fabricates. When it does not know, it says so. Retrieval grounds every answer.

2. Problem Statement — Deep Analysis

Problem 1: Cloud Dependency

Every major AI system today requires stable internet, low-latency connections, and continuous connectivity. For 340 million Indians in areas with unreliable connectivity, this means zero AI access during the moments they need it most — emergencies, remote work, travel.

Impact

340M users structurally excluded. Data sent to foreign servers. ₹1,660/month unaffordable.

Problem 2: Language Exclusion

Every major AI system is fundamentally English-first. GPT-4 was trained on data that is ~92% English. When these systems 'support' Hindi or Tamil, quality degrades sharply for complex reasoning, code-switching (Hinglish), and cultural context.

Impact

Hindi: 530M speakers underserved. Bengali: 100M. Tamil: 80M. Telugu: 80M. Marathi: 83M.

Problem 3: Hallucination in High-Stakes Contexts

LLMs generate text by predicting likely next tokens. Without retrieved context, they fabricate plausible-sounding but incorrect information. For casual use, inconvenient. For medical guidance, legal rights, crop decisions — dangerous.

Impact

Offline systems have no live grounding fallback. Wrong medical advice. Incorrect legal guidance. Bad agricultural decisions.

Problem 4: Stateless AI — No Memory

Every conversation with ChatGPT, Gemini, or Claude starts from zero. No memory of who you are, what you discussed before, or what you are working on. Each session is a blank slate. When AI systems try to solve this by storing full history, they face RAM explosion on mobile.

Impact

Repetitive interactions. Shallow personalization. KV cache bloat crashes mobile apps.

Problem 5: Hardware Inefficiency

Modern AI assumes 80GB VRAM, 400W TDP, and cloud GPUs. Real users have 4GB RAM, 5W thermal budgets, and battery constraints. Naive mobile AI ports result in overheating, 20-minute battery drain, and OS process termination.

Impact

Local AI becomes unusable on 95% of the world's smartphones.

Problem 6: The Cost Barrier

ChatGPT Plus, Gemini Advanced, and Claude Pro each cost \$20/month (₹1,660). The median rural Indian monthly income is ₹8,000–12,000. AI subscriptions represent 14–21% of income — structurally impossible.

Impact

The people who could benefit most from AI are precisely those who cannot afford it.

3. Existing Solutions Comparison

Feature Comparison Matrix

Dimension	ChatGPT	Gemini	Apple Intelligence	Samsung AI	ISHA AI
Offline capability	None	None	Partial	Partial	Full
Privacy (no data leak)	No	No	Partial	No	Yes — absolute
Indian language quality	Poor	Moderate	Poor	Moderate	Native
Works on ₹15k phone	No	No	No (iOS)	No	Yes
Free forever	No	No	HW cost	Weak	Yes — full
Live camera (offline)	Cloud only	Cloud only	Limited	Limited	On-device
Voice (offline)	Cloud only	Cloud only	Cloud only	Limited	Fully offline
Persistent memory	Limited	Limited	Limited	None	Semantic graph
India domain knowledge	None	None	None	None	Full packs
Peer mesh intelligence	None	None	None	None	Unique — ISHA only
Works during disasters	No	No	No	No	Yes
Monthly cost	₹1,660	₹1,660	₹0 (HW)	₹0 (weak)	₹0 — forever

Why Each Competitor Fails ISHA's Market

ChatGPT

Aspect	Detail
Strengths	Best general reasoning quality, massive training data, strong tool use, large developer ecosystem.
Critical Weaknesses	Zero offline capability. Hindi/regional quality significantly below English. No knowledge of Indian schemes. All conversations on OpenAI servers. ₹1,660/month for Plus tier.
Engineering Limitation	GPT-4o requires ~80GB VRAM just for weights. Porting to mobile is physically impossible with current hardware.

Google Gemini

Aspect	Detail
Strengths	Gemini Live has impressive real-time video (cloud). Gemini Nano demonstrates on-device viability. Strong multilingual training data.
Critical Weaknesses	Gemini Nano only on Pixel phones (~8% Android market). Still requires cloud for complex reasoning. Google's business model is data-driven advertising.
Engineering Limitation	Gemini Nano (~1.8B params) cannot match ISHA's Tier 1+2 combination with retrieval augmentation.

Apple Intelligence

Aspect	Detail
Strengths	Best hardware-software integration. Private Cloud Compute is genuinely privacy-preserving. Deep iOS integration.
Critical Weaknesses	iOS exclusive — excludes 75% of India's smartphone market. iPhone costs ₹50,000+ minimum. No Indian domain knowledge. No user-customizable knowledge base.
Engineering Limitation	Structurally inaccessible to India's mass market. Cannot run on Android. Requires Apple Silicon.

4. Full System Architecture

High-Level Architecture

ISHA Runtime Stack (Top to Bottom)		
USER INTERFACE LAYER		
Voice Interface Text Interface Camera/Vision Haptic		
ISHA INTERACTION LAYER		
Wake Word Engine STT Pipeline TTS Engine Vision Pipeline		
COGNITIVE RUNTIME ENGINE		
Intent Classifier Context Builder Adaptive Model Router		
└─ ORCHESTRATION CORE		
INFERENCE ENGINE	RETRIEVAL ENGINE	MEMORY ENGINE
T0: 350M param	FAISS Index	L1: Hot RAM
T1: 1.5B param	Knowledge Packs	L2: Warm SQLite
T2: 2B param	Graph DB	L3: Cold vectors
T3: 7B param*	Personal memory	L4: Encrypted
VERIFICATION PIPELINE		
Grounding Check Confidence Score Offline Boundary Gate		
DEVICE COGNITION LAYER		
Battery Thermals Apps Calendar Permissions		
* T3 only on 8GB+ RAM devices		

Complete Request Lifecycle — 20 Steps

Step	Action	Component	Time
1	User says 'Hey ISHA'	DSP Wake Word	<50ms
2	Wake phrase confirmed	CPU Verification Model	<300ms total
3	Main process resumes from hibernate	OS / ISHA Runtime	<300ms
4	STT pipeline activates	Whisper.cpp	Immediate
5	User speaks query	Whisper streaming	Real-time
6	Transcription completes (VAD)	Whisper VAD	~200ms after speech
7	Text enters Cognitive Runtime	Orchestration Core	Immediate

Step	Action	Component	Time
8	Intent Classifier analyzes query	T0 Nano Brain	<200ms
9	Context Builder assembles memory + device state	Memory Engine	<50ms
10	Retrieval Engine queries FAISS + Packs	FAISS + Cross-encoder	<60ms
11	Retrieved context assembled (2–4KB)	Retrieval Orchestrator	<10ms
12	Model Router decides T0/T1/T2/T3	Adaptive Model Router	<10ms
13	Selected model generates response (streamed)	llama.cpp / CoreML	500ms–8s
14	Verification Pipeline checks response	Grounding + Confidence	<100ms
15	TTS converts response to speech	Kokoro-82M	<150ms first chunk
16	Response played to user	Audio system	Streaming
17	Semantic anchors extracted	T0 (async)	Background
18	Anchors written to L2 memory	Memory Engine	Background
19	T2/T3 model unloads (if used)	Adaptive Model Loader	Immediate
20	System hibernates, DSP resumes	Runtime + DSP	<500ms

5. Model Architecture & AI Stack

The Four-Tier Model System Overview

Tier	Name	Model	Disk Size	RAM	Tok/s	Scope
T0	Nano Brain	Phi-3 Mini pruned → 350M	~210MB	~180MB	20–25	Always resident — routing, quick answers
T1	Smart Brain	Qwen2.5-1.5B-Instruct Q4_K_M	~950MB	~680MB	15–18	Primary workhorse — Hindi/multilingual
T2	Deep Brain	Gemma-3-2B-IT Q4_0	~1.4GB	~1.1GB	10–14	Burst reasoning — complex tasks (90s limit)
T3	Ultra Brain	Mistral-7B-Instruct Q4_K_M	~4.1GB	~4.2GB	4–7	Flagship only — maximum intelligence

Why Qwen2.5-1.5B for Tier 1 — The Critical Decision

This is the single most important model selection in the entire ISHA stack. The choice directly determines quality for 800 million Indian language speakers.

Model	Hindi Reasoning	Multilingual MMLU	English MMLU	Disk Size (Q4)
Qwen2.5-1.5B (CHOSEN)	Best in class	Best in class	Strong	~950MB
Gemma-2-2B	Moderate	Moderate	Strong	~1.4GB
Phi-3.5-mini	Weak	Weak	Best in class	~2.2GB
TinyLlama-1.1B	Very weak	Very weak	Weak	~650MB
Llama-3.2-1B	Weak	Moderate	Moderate	~700MB

Qwen2.5 was trained by Alibaba on a massive multilingual corpus with particular strength in Indian language reasoning. At 1.5B parameters it outperforms alternatives 2–3x larger on Hindi benchmarks.

Inference Engines

Platform	Engine	Why Chosen	Acceleration
Android (all tiers)	llama.cpp	Built specifically for on-device LLM inference. Best GGUF quantization support. ARM NEON optimized. Fine-grained memory control.	NNAPI backend + Vulkan GPU + ARM NEON CPU
iOS T0/T1	CoreML (ANE primary)	Apple Neural Engine delivers 38 TOPS at fraction of GPU power. ANE decode phase dramatically reduces battery drain.	ANE decode + GPU prefill (split-phase)
iOS T2/T3	CoreML (GPU prefill)	Heavy prefill phase benefits from parallel GPU. Token decode benefits from efficient ANE sequential execution.	GPU prefill + ANE decode

Platform	Engine	Why Chosen	Acceleration
	+ ANE decode)		

Embedding & Retrieval Models

Model	Size	RAM	Purpose	Why Chosen
all-MiniLM-L6-v2	~22MB	~28MB	Query/document embedding for semantic search	22MB delivers embedding quality rivaling models 10x larger. Always resident with T0.
ms-marco-MiniLM-L-6- v2	~45MB	~52MB	Reranking retrieved results for precision	Cross-encoder reranking significantly improves retrieval accuracy — critical for grounding small models.

6. Voice & Wake Word System

Two-Stage Wake Detection Architecture

Always-listening voice systems face a fundamental tension: responsive (must hear 'Hey ISHA' instantly) vs battery-preserving (cannot run a large model continuously). ISHA solves this with a two-stage system.

Stage	Runs On	Model Size	Power	Accuracy	Role
Stage 1: DSP Keyword Spotter	Dedicated DSP chip (Qualcomm Hexagon / Apple Always-On DSP)	2MB	~1mAh/hour	~85% (intentionally loose)	Always listening. Wakes CPU on probable match.
Stage 2: CPU Verification	CPU (active <200ms only)	8MB	Negligible (milliseconds)	>99% on Stage 1 candidates	Confirms or rejects Stage 1 trigger. Full ISHA activates only on confirmation.

Battery math: DSP spotter uses ~24mAh/day = ~0.73% of iPhone 15 Pro battery. Equivalent to having Bluetooth on. Acceptable.

Speech-to-Text — Whisper.cpp

Tier	Model	Size	WER Hindi	WER English	Speed
ISHA Lite	whisper tiny	75MB	18.2%	8.1%	0.8x realtime
ISHA Standard	whisper small	244MB	12.4%	5.3%	1.4x realtime
ISHA Pro	whisper medium	769MB	9.1%	4.1%	3.2x realtime

Whisper was trained on 680,000 hours of multilingual audio including significant Indian language content. It handles Hinglish (Hindi-English code-switching) natively — critical for how hundreds of millions of Indians actually speak.

Text-to-Speech — Kokoro-82M


Voice	Language	Style	Size
ISHA	English	Neutral, warm, professional	Included in 82MB base
ISHA Hindi	Hindi-first	Warm, conversational	Included in 82MB base
Regional packs	Tamil, Telugu, Bengali, Marathi,	Native prosody per language	~15MB each (optional download)

Voice	Language	Style	Size
	Gujarati, Kannada, Malayalam, Punjabi		

7. Vision & Live Camera Architecture

Sparse Event-Driven Vision — The Core Design

Vision AI on mobile has the hardest constraints of any ISHA feature. A VLM at useful quality starts at 1.7B parameters. Continuous camera processing causes device overheating in 3–5 minutes. ISHA's answer: analyze intelligently — only when something changes, only in the changed region, only when thermals permit.



Vision Capabilities by Tier

Capability	Lite	Standard	Pro
Photo analysis	Yes	Yes	Yes
OCR (text in image)	Yes	Yes	Yes
Document scan + extract	Yes	Yes	Yes
Bill / receipt reading	Yes	Yes	Yes
Medicine label reading	Yes	Yes	Yes
Plant / crop disease (photo)	Yes	Yes	Yes
Live camera scene	No	Yes	Yes
Real-time text translation	No	Yes	Yes
Complex visual reasoning	No	Basic	Full

Capability	Lite	Standard	Pro
Handwriting recognition	Basic	Full	Full
Video understanding	No	No	Sparse 1fps

8. Knowledge Pack System

Analogy: ISHA is an open-book exam student. Knowledge packs are perfectly organized notes. A smaller memory (model) + perfect notes beats a massive memory (huge model) + no notes.

Official Pack Catalog

Pack	Contents	Installed Size	Query Latency
Bharat Core (Free, Always)	NCERT Class 1–12, Indian Constitution, 500+ govt schemes, Essential medicines, Emergency first aid, Legal rights basics	~140MB	<15ms
Exam Warrior	UPSC Prelims+Mains (15 yrs), JEE Main+Advanced (10 yrs), NEET, SSC CGL+CHSL, State PSC (12 states), Banking exams	~280MB	<20ms
Kisan Saathi	Crop disease DB (2000+ diseases), Pest ID+control, Crop calendar, Soil health, Govt subsidies, Mandi price history	~180MB	<15ms
Swasthya (Health)	Symptom checker, Medicine interactions, First aid, Mental health (Hindi), Maternal+child health, Nutrition guide	~210MB	<18ms
Legal Sahayak	Consumer rights, RTI procedure, Labor laws, Domestic violence support, Court procedures	~120MB	<12ms
Bhasha Pack	Regional language enhancement for 8 Indian languages, Dialect-specific knowledge, Cultural context	~90MB	<12ms

Pack Architecture — Internal Structure

Pack File Layout	
kisan_pack/	
├─ manifest.json	# Pack metadata, version, Ed25519 signature
├─ index.faiss	# Pre-built vector index (fast search)
├─ chunks/	# Pre-processed text chunks (256 tokens each)
│ └─ wheat_diseases.bin	
│ └─ rice_cultivation.bin	
│ └─ ... (thousands of chunks)	
├─ embeddings/	# Pre-computed float16 vector embeddings
│ └─ vectors.bin	# 384 dimensions per chunk
├─ graph/	# Knowledge graph entity relationships
│ └─ entities.db	# SQLite: entity → related entities
└─ metadata/	# Source attribution per chunk
└─ sources.json	

Install-Time Pre-Processing

All heavy work happens ONCE at install time, never at query time. This is why ISHA achieves <25ms retrieval on mobile.

Step	Operation	Output	Time
1	Document chunking (256-token chunks, 32-token overlap)	Hundreds of chunks per pack	~30s
2	Embedding generation (MiniLM-L6 on each chunk)	384-dim float16 vectors	~2–4 min
3	FAISS index building	.faiss binary index file	~30s
4	Knowledge graph construction (entity extraction)	SQLite entity graph	~1 min
5	Pack ready	All future queries: <25ms retrieval	Done once forever

Community & Enterprise Packs

Pack Type	Creator	Use Case	Revenue
Official Packs	ISHA Team	National curriculum, health, legal, farming	Free always
Community Packs	Any teacher, doctor, expert	Class-specific, regional, domain packs	Free — creator attribution
Enterprise Packs	Companies, hospitals, universities	Internal protocols, private knowledge bases	B2B licensing

9. Memory Architecture

Four-Layer Memory Hierarchy

Layer	Technology	What Lives Here	Max Size	Lifetime	Encryption
L1: Hot	RAM (native heap)	Active KV cache, current turn context only	50KB strict	Destroyed after each model response	Process isolation (RAM)
L2: Warm	SQLCipher (encrypted SQLite)	Session semantic anchors, user preferences	200KB session + 10MB prefs	24 hours → auto-delete	AES-256-GCM, device keystore
L3: Cold	Compressed float16 vectors + SQLite	Long-term behavioral + knowledge profile	50MB	Permanent until user deletes	AES-256-GCM, Secure Enclave (iOS) / Keystore HW (Android)
L4: Cognitive Fingerprint	JSON document (encrypted)	Who the user is — structured profile	500KB	Permanent, user-visible and editable	Biometric-bound — Face ID / fingerprint required

Semantic Compression — The Core Innovation

Instead of storing raw conversation text, ISHA extracts and stores only semantic meaning. Compression ratio: 80x. A 32KB conversation becomes a 400-byte semantic anchor.

Aspect	Traditional AI (Full History)	ISHA (Semantic Compression)
Storage per conversation	32KB–500KB raw tokens	~400 bytes semantic anchors
RAM at query time	Grows unboundedly with history	Bounded — max 50KB hot cache
Privacy risk	High — full conversation stored	Low — only meaning, no raw text
Long-term utility	Declines (too much noise)	Improves (signal only)
Context reconstruction	Exact replay	Approximate — from semantic anchors

24-Hour Chat Lifecycle

Chat Lifecycle Flow
SESSION STARTS → Ephemeral session created in RAM only
CONVERSATION PROCEEDS
Every 3 exchanges: Semantic compression runs (T0, background)
Anchors written to L2 SQLite (encrypted)
Raw conversation stays in RAM only — never written to disk

23 HOURS → ISHA notification:

'Your conversation from yesterday will be deleted in 1 hour.'

[Save to device] – encrypted, warns about storage impact

[Download] – exports to Files app, ISHA copy deleted

[Delete now] – gone immediately, cryptographically erased

[Remind later] – 30 more minutes

24 HOURS → Auto-delete if no action taken

L2 cleared | L4 cognitive fingerprint updated from anchors

Only the 400-byte semantic summary survives.

Raw conversation: permanently gone.

10. Retrieval Engine (RAG)

The Retrieval Pipeline

Step	Operation	Time	Detail
1	Query preprocessing	<5ms	Language detection, intent type, entity extraction
2	Query embedding (MiniLM-L6)	~8ms	Query → 384-dim vector
3	FAISS vector search	~12ms	Top-20 candidate chunks from active packs
4	Cross-encoder reranking	~35ms	ms-marco reranker scores each candidate → top 5 selected
5	Knowledge graph augmentation	<5ms	Related entity traversal → additional context chunks
6	Context assembly	<5ms	System prompt + user profile + 5 passages (~500 tokens) + query
7	Model receives grounded context	—	Hallucination rate: <5% vs ~40% without RAG for small models

Online Mode — MCP + Live RAG

When internet is enabled, ISHA adds live data sources to retrieval, making it potentially better than ChatGPT: personal knowledge + live internet + domain packs — all merged locally.

MCP Tool	Data Source	Use Case
Web Search	Brave Search / Tavily	Current events, live research, real-time information
Wikipedia	Wikipedia API	Encyclopedic knowledge, recent article updates
Weather	Weather APIs	Forecast data for Kisan pack, travel planning
News	News feed APIs	Current affairs, exam notifications
Device Calendar	Local (no MCP)	Appointment-aware responses
Custom MCPs	User-configured	GitHub, Google Drive, company tools, etc.

Privacy guarantee in online mode: PII stripped from all MCP calls before transmission. User identity, location, and conversation history never leave the device.

11. ISHA Mesh — Peer Intelligence Network

ISHA's most unique innovation. No other production AI system distributes inference across nearby consumer devices. This is ISHA's moat.

The Concept

ISHA Mesh allows multiple ISHA-enabled devices on the same local network or Bluetooth range to share inference load. A user with a ₹15,000 phone accesses reasoning quality far beyond their device's capability by leveraging nearby idle compute.

Scenario	Without Mesh	With ISHA Mesh
Child on Redmi C65 (₹15k) asks complex JEE question	T1 only — limited reasoning quality	3 family phones collaborate — equivalent to 4.5B model quality
30-student classroom	Each student limited to own device tier	Collective compute ~30x — teacher's T2 available to all
Village with 50 ISHA phones on WiFi	Individual Lite capability	Collective intelligence rivals premium cloud AI — fully offline

Mesh Protocol Stack

Layer	Technology	Details
Discovery	mDNS (_isha-mesh._tcp.local)	Broadcasts: device tier, available capacity, battery%, thermal state
Transport	Local TCP socket	AES-256-GCM encrypted per session. ECDH key exchange.
Coordination	Gossip protocol	Dynamic coordinator election — highest capability + best battery
Task distribution	Sub-query decomposition (T0)	Parallelizable queries split by sub-topic, matched to node capability
Privacy	Differential privacy design	Sub-queries stripped of personal context. No raw user data shared.

Privacy in the Mesh — Critical Design

Category	Shared Across Mesh	Never Shared
Query data	Sub-query text (impersonal)	User identity, personal history, full conversation
Results	Sub-response text	Nothing from worker nodes retained after sending
Identity	Hashed device ID (not raw)	Real device ID, IMEI, phone number
Media	Never	Voice audio, images, video — always stay on originating device

12. Security & Privacy Deep Dive

Threat Model

Threat Actor	Goal	ISHA Defense
Malicious apps on device	Steal conversation data	Process isolation + AES-256-GCM encrypted storage
Physical device theft	Access user data	Biometric-bound encryption — Secure Enclave / Android Keystore HW
ISHA company itself	Access user data	Architecturally impossible — we hold no user data
Government legal requests	Compel data handover	We have no server, no user data — nothing to hand over
Network eavesdropping	Intercept MCP queries	All MCP calls via HTTPS/TLS. No queries in offline mode.
Mesh eavesdropping	Intercept mesh traffic	AES-256-GCM per session, ECDH forward secrecy
Malicious packs	Inject harmful content	Ed25519 cryptographic signature verification on all packs

What ISHA Cannot Do — By Architecture

Capability	Why It Cannot Happen
Send user conversations to any server	No code path exists for this in offline mode. Architecturally impossible.
Access data without biometrics	Secure Enclave / Android Keystore HW prevents access without Face ID / fingerprint.
Be compelled to hand over user data	We hold no user data on any server. Nothing exists to compel.
Run hidden tracking or analytics	All telemetry is opt-in, on-device only, never transmitted without consent.
Update silently to add data collection	App Store / Play Store review required + user consent for all updates.

13. Performance Engineering

Latency Budget

Stage	Target	Maximum Acceptable	Key Optimization
Wake word detection	<300ms	500ms	DSP chip — runs always, zero CPU cost
STT transcription start	<200ms	400ms	Whisper streaming — first tokens as user speaks
Intent classification (T0)	<200ms	400ms	T0 always resident — no load time
Retrieval (FAISS + rerank)	<60ms	150ms	FAISS index in RAM — no disk I/O at query time
T1 first token	<500ms	1000ms	Predictive pre-loading during STT phase
T1 full response (100 tokens)	<6s	10s	Token streaming — user reads while model generates
TTS first audio chunk	<150ms	300ms	Kokoro streaming — audio plays before full response done
Total: voice to first voice	<8s	15s	Parallel processing of all stages

Thermal State Machine

Thermal State	Device Temp	Capabilities	User Message
NOMINAL	< 35°C	All tiers, full context (512 tokens), vision enabled, mesh participation	Normal operation
WARM	35–38°C	T2 disabled, T1 max, reduced context (256 tokens), photos only	Transparent — no message
HOT	38–41°C	T1 disabled, T0 only, retrieval-first answers, vision disabled	'Running warm — simple answers only until device cools.'
CRITICAL	> 41°C	All inference suspended, retrieval-only	'Device needs to cool. Here is what I found in your notes:'

14. End-to-End User Journey

Ramesh — Wheat Farmer, Vidarbha, 7:15 AM

Device: Redmi Note 13, 6GB RAM, ISHA Standard. No internet. Standing in field.

Time	Event	Component	Detail
7:15:00.000	Ramesh says 'Hey ISHA'	DSP Keyword Spotter	Wake phrase detected
7:15:00.180	Wake phrase confirmed	CPU Verification Model	98.3% confidence
7:15:00.280	ISHA resumes from hibernate	Runtime	<300ms total wake time
7:15:00.320	ISHA plays earcon chime	Audio	'I am listening'
7:15:02.100	Ramesh speaks in Marathi about yellow wheat leaves + shows camera	Whisper STT + Vision	Parallel audio transcription + camera frame captured
7:15:06.600	Transcription complete	Whisper VAD	Marathi text produced
7:15:06.640	Parallel processing starts	Orchestration Core	Intent (T0) + Vision (VLM) + Retrieval (FAISS) run simultaneously
7:15:07.400	T1 loaded	Adaptive Model Loader	Pre-loaded during STT phase — fast
7:15:07.600	Context assembled	Retrieval Orchestrator	Ramesh profile + visual analysis + Kisan Pack chunks + query
7:15:07.700	T1 inference begins	Qwen2.5-1.5B via llama.cpp	Grounded context provided
7:15:08.100	First token generated	T1	~400ms from context assembly
7:15:08.200	ISHA speaks in Marathi	Kokoro TTS	Streaming audio — Ramesh hears response as it generates
7:15:14.000	Full response delivered	Complete pipeline	Diagnosis + treatment + govt subsidy info
7:15:14.200	Semantic anchor written to L2	Memory Engine (async)	'wheat yellowing, possible N/Zn deficiency, treatment given'
7:15:14.400	T1 unloads	Adaptive Model Loader	~680MB returned to system
7:15:14.500	ISHA hibernates, DSP resumes	Runtime + DSP	Battery impact of full interaction: ~0.8%

Total interaction: 12.5 seconds. Agricultural guidance that would cost ₹500 from a consultant. Zero internet. Zero cost. In Ramesh's language.

15. Technical Tradeoff Analysis

Decision	Chosen Approach	Why Chosen	What Was Sacrificed
Model size vs quality	Small model (1.5B) + strong RAG	Supports 4GB devices (80% of market). Retrieval compensates for size. Hallucination <5%.	Open-domain queries without relevant packs are weaker.
Memory: ephemeral vs persistent	Ephemeral (24hr) + semantic extraction	Privacy. Bounded storage. No RAM explosion. User trust.	Exact conversation recall unavailable unless user saves.
Inference engine: llama.cpp vs ONNX	llama.cpp for Android	Built for LLMs. Best GGUF support. Fine-grained mobile memory control.	Slightly more integration complexity vs ONNX.
Vector store: FAISS vs ChromaDB	FAISS	C++ library — 5MB, <5ms queries on mobile CPU.	ChromaDB/Qdrant have better developer tooling — unsuitable for mobile.
Wake word: DSP vs always-on CPU	DSP two-stage	~0.73% battery/day vs 8–15% for always-on CPU.	Slightly higher false rejection in loud environments.
KV cache: full vs sliding window	Sliding 512-token window + semantic checkpoints	Bounded RAM (~50KB vs 400MB for 7B full context).	Deep context reconstruction is approximate, not exact.

16. Database Architecture

Storage Layer Overview

Layer	Technology	Contents	Max Size	Lifetime
In-Process RAM	Native heap	Active KV cache, current session context	50KB	Current inference turn only
SQLCipher SQLite	Encrypted SQLite (AES-256-GCM)	Sessions, semantic anchors, preferences, pack index, query cache	<10MB	24hr sessions + permanent prefs
FAISS Binary Indexes	Meta FAISS + custom binary	Knowledge pack vector indexes	10–300MB per pack	Until pack uninstalled
Vector Store	Custom float16 + SQLite metadata	Long-term user memory graph (L3/L4)	<50MB	Permanent until user deletes
Model Files	Memory-mapped GGUF / CoreML	Model weights (T0–T3)	200MB–4.1GB	Until model uninstalled

Key SQLite Tables

Table	Primary Key	Key Columns	Purpose
sessions	session_id (TEXT)	created_at, expires_at, status, language, interaction_count	Manages 24hr chat lifecycle
semantic_anchors	anchor_id (TEXT)	session_id, anchor_type, content (JSON), confidence, promoted	Stores compressed session knowledge
user_preferences	key (TEXT)	value, updated_at	Permanent user settings (language, voice, tier preference)
installed_packs	pack_id (TEXT)	name, version, index_path, chunks_path, signature, is_active	Manages knowledge pack installations
query_cache	query_hash (SHA-256)	response, pack_versions, created_at, hit_count	Anonymous query result cache
performance_log	log_id (TEXT)	tier_used, inference_ms, peak_ram_mb, device_temp_c, battery drain	On-device performance telemetry (never sent)

17. Platform Architecture — iOS vs Android

Dimension	Android Lite (₹10k–20k)	Android Standard (₹30k–60k)	iOS Pro (iPhone 12+)
RAM budget (ISHA)	~1.0GB max	~2.0GB max	~2.8GB max
Inference engine	llama.cpp + NNAPI partial	llama.cpp + NNAPI full + Vulkan	CoreML + ANE native
Tiers available	T0 + T1 only	T0 + T1 + T2	T0 + T1 + T2 + T3
Vision	Static images only (MobileVLM-1.7B)	Live sparse video (MobileVLM-3B)	Full Gemini-like (LLaVA CoreML)
TTS quality	Basic eSpeak	Neural Piper	Neural Kokoro + AVSpeechSynthesizer
Wake word	DSP (Qualcomm Hexagon)	DSP (Qualcomm Hexagon)	Always-On DSP (Apple)
Secure storage	Android Keystore (basic)	Android Keystore (full HW)	Secure Enclave + Face ID binding
Background tasks	JobScheduler	WorkManager	BGProcessingTask
OS integration	None	Basic Android Intents	AppIntents + Siri Domains + Live Activities

What is IDENTICAL across all platforms: Privacy architecture, ephemeral memory design, retrieval-first philosophy, knowledge packs, offline-first guarantee, ISHA Mesh protocol, hallucination verification, core T0 behavior.

18. Hardware Tiering System

Device Capability Detection

At first launch, ISHA runs a 20-second hardware profiler to determine exactly what the device can support. This generates a permanent capability manifest that governs all future routing decisions.

Device Category	Example Devices	RAM	ISHA Tier	Key Unlocks
Budget Android	Redmi 13C, Moto G34, Realme C65	4GB	ISHA Lite	T0+T1, Hindi/English, Core pack, basic memory, text+voice
Mid Android	Redmi Note 13, Samsung A35, Moto G84	6–8GB	ISHA Standard	T0+T1+T2, all 10 languages, full packs, live vision (sparse)
Flagship Android	Samsung S24, OnePlus 12, Pixel 8 Pro	8–12GB	ISHA Standard+	T0+T1+T2+T3 (8GB+), deeper OS integration
iPhone (A14/A15)	iPhone 12, 13 series	6GB	ISHA Pro	CoreML/ANE, T0+T1+T2, full iOS integration, Siri handoff
iPhone Pro (A16+)	iPhone 14 Pro, 15 Pro, 16 series	8GB	ISHA Pro Max	T3, full vision, ambient intelligence, Deep OS integration
iPad M-series	iPad Pro M2/M3/M4	8–16GB	ISHA Ultra	All tiers, video understanding, full agent capabilities

19. Infrastructure & Deployment

Server Infrastructure — Minimal by Design

ISHA's server infrastructure is 99.5% smaller than equivalent cloud AI because all inference runs on-device. Monthly server costs for 1 million users: ~\$300. Compare to cloud AI: \$50,000–200,000/month.

Component	Technology	Purpose	Monthly Cost (1M users)
Pack CDN	Cloudflare R2	Knowledge pack distribution — static file serving	~\$50–200
Pack Store API	Lightweight REST (Go/FastAPI)	Pack browsing, metadata, submission endpoint	~\$50
Pack Signing Service	Ed25519 signing server	Cryptographic signature generation for verified packs	~\$20
Crash Reporting	Sentry (self-hosted)	Anonymous stack traces only — no user data	~\$30
Total	—	Complete ISHA backend for 1M users	~\$150–300/month

Cost Comparison

Service	Architecture	Cost for 1M users / 5 queries/day	Revenue needed to break even
ChatGPT (cloud)	GPU server inference	~\$750,000/month	37,500 paying users at \$20/month (infrastructure only)
ISHA (on-device)	Device inference + minimal CDN	~\$300/month	~15 enterprise pack customers

20. Testing Strategy

Test Category	Method	Gate Criteria	Frequency
Model quality (Hindi/regional)	Automated benchmark: 500+ test cases per language	Must hit >95% of float16 baseline on Hindi benchmark	Every model update
Latency benchmarks	ADB-connected device farm, 12 device models	No >20% regression from baseline on any P95 metric	Every build
Privacy: no network calls	Network traffic monitor + 100 offline queries	ZERO outbound requests — any failure BLOCKS release	Every build
Memory pressure (iOS)	XCUITest simulated memory warning during T2 load	T2 unloads in <2s, T0 remains operational, no data loss	Every build
Android LMK survival	Fill device RAM until pressure, verify ISHA behavior	Process priority correct, state recovers from L2 after kill	Every build
24-hour auto-delete	Mock time to +25 hours, query L2 SQLite	Zero rows for expired sessions	Every build
Thermal behavior	30-min stress test at 25°C ambient	Device temp <42°C, no thermal throttle, <18% battery drain	Weekly
Language quality (human eval)	Panel: 5 Hindi + 3 Tamil + 3 Telugu + 3 Bengali + 2 Marathi speakers	No language scores below 3.5/5 on grammar/accuracy/naturalness	Weekly in dev, monthly post-launch

21. Engineering Mentor Review

What Is Excellent

Strength	Why It Matters
Retrieval-first philosophy	Industry is slowly learning what ISHA bakes in from day one: small model + great retrieval beats large model + no retrieval. Google, Apple, Microsoft all validating this direction in 2025–2026.
ISHA Mesh peer intelligence	No production AI system distributes inference across nearby consumer devices. Real competitive moat requiring years to replicate — not because secret, but because trust model + discovery + synthesis require careful engineering.
Modular knowledge packs	Fundamentally changes who benefits from AI. Teacher creating a class pack is a natural behavior in India (WhatsApp groups share study materials constantly). ISHA packs are the AI-native equivalent.
Privacy by architecture	Designing so ISHA company cannot access user data — not just choosing not to — is the correct approach. No legal jurisdiction can compel handover of data that does not exist on any server.
Event-driven cognition lifecycle	Wake → Reason → Compress → Unload → Hibernate aligns AI lifecycle with OS scheduling principles. Most ML engineers miss this systems insight entirely.

What Needs Work — Honest Assessment

Weakness	Severity	Issue	Recommendation
T3 (7B) on mobile RAM	High	4.2GB on 8GB device leaves <3.8GB for OS + apps + T0. Android often has only 4–5GB available after ROM.	Design T3 as best-effort. Aggressive thermal + memory gates. Market as 'on capable devices' without hard promises.
90-second T2 limit UX	Medium	Complex analysis (10-page legal doc, multi-step physics) may need >90 seconds. Cutting off mid-task is terrible UX.	Checkpoint-based inference: T2 saves thinking state every 30s. Partial results presented with 'ISHA needs a moment to continue...'
Voice latency on Lite	Medium	T1 at 100 tokens on Cortex-A55 may take 8–12 seconds. Exceeds 8-second target.	For Lite voice: 50-token response limit, or text display while voice generates in parallel.
Hinglish TTS prosody	Medium	Kokoro produces good Hindi and English separately. Mid-sentence code-switching prosody is poor.	Sentence-level language detector. Hindi phonemes for Hindi sentences, English for English. Concatenate.
Community pack moderation at scale	High	Manual review works at 1,000 packs. Impossible at 10,000+.	Automated content scanning pipeline using T2 on ISHA's server. Human review queue for flagged content.

Hidden Risks

Risk	Severity	Description	Mitigation
Qwen2.5 license at scale	High	>100M monthly active users requires separate agreement with Alibaba.	Start negotiation at 10M users. Evaluate alternative models proactively.
NNAPI fragmentation	Medium	NNAPI implementation quality varies dramatically by chip OEM. May be slower than CPU on some devices.	Maintain device capability database. Auto-benchmark NNAPI vs CPU at first launch. Fall back when NNAPI is slower.
iOS background execution	Low–Medium	BGProcessingTask requires charging device. Memory distillation may not complete.	Opportunistic foreground distillation when app is briefly opened. Background is secondary.
Pack freshness	Medium	Packs go stale. Government schemes change, exam patterns update.	Clear version dates on all packs. Expiry warnings. Differential updates (changed chunks only).

22. Production Readiness Assessment

Dimension	Score	Rationale
Core AI Architecture	9/10	Retrieval-first + tiered inference is production-grade. Industry-validated direction.
Privacy & Security	9/10	Secure Enclave + no-server design is best-in-class for consumer AI.
Scalability (user growth)	8/10	No server to scale — scales with device hardware. Pack CDN trivially scales.
Offline reliability	9/10	True offline-first with no hidden dependencies. Tested against network kill.
Multilingual quality	8/10	Qwen2.5 + Whisper strong. Hinglish TTS needs dedicated model.
Performance — Lite devices	6/10	Latency on Cortex-A55 will be tight. Needs optimization work.
Performance — Pro devices	9/10	ANE + CoreML split-phase pipeline is production-ready. Excellent.
Knowledge pack ecosystem	7/10	Architecture solid. Community pack moderation at scale is TBD.
Voice pipeline	7/10	Whisper + Kokoro good. Hinglish TTS weak. Voice latency on Lite tight.
Vision pipeline	7/10	Sparse design is clever and production-safe. Lite accuracy limited.
ISHA Mesh	6/10	Novel architecture. Real-world latency and reliability needs field testing.
OVERALL	7.8/10	Exceptional design for constraints. Primary risks are engineering challenges, not architectural flaws.

23. Product Roadmap

Phase	Timeline	Focus	User Target
Phase 1: Foundation	Months 1–6	T0+T1 Android, text interface, Hindi+English, Bharat Core pack, basic RAG, 24hr memory, Play Store alpha	1,000 beta users
Phase 2: Voice & Language	Months 7–12	Voice pipeline (Whisper+Kokoro), 6 Indian languages, T2 on Standard, Exam+Health packs, iOS beta	100,000 users
Phase 3: Vision & Mesh	Months 13–18	Image understanding all tiers, live camera Standard+Pro, ISHA Mesh v1, Kisan pack visual disease ID, community packs	1 million users
Phase 4: Platform	Months 19–24	Cross-platform mesh, MCP online mode, OEM partnerships (Micromax/Lava pre-install), Enterprise packs, T3 flagship	10 million users
Phase 5: Infrastructure	Months 25–36	ISHA OS layer (deep Android integration), feature phones (smaller models), satellite connectivity, classroom mesh	100 million users

24. Failure Scenarios & Recovery

Failure	Detection Method	Recovery Steps	User Experience
Model crash mid-inference (OOM)	Exception handler / truncated output detector	1. Catch exception 2. Unload all tiers 3. T0 reloads (always resident) 4. Offer retry with T1 instead of T2 5. Log to on-device perf log	'I ran into a memory issue. Let me try a simpler approach.'
FAISS index corruption	SHA-256 checksum on launch + FAISS load validation	1. Mark pack corrupted in SQLite 2. Notify user 3. Offer re-download or P2P mesh transfer 4. Continue without that pack	'Kisan Pack needs reinstalling. Other features still work.'
ISHA Mesh node timeout	30-second coordinator timeout per sub-task	1. Coordinator handles missing sub-tasks locally 2. Merge available + local results 3. Remove unresponsive node (60s TTL) 4. Node auto-rejoins when available	Slightly longer response. No visible failure shown to user.
iOS background task killed	iOS expiration handler callback	1. Save partial distillation state 2. Mark session 'distillation_incomplete' 3. Resume on next foreground open 4. If session expires: distill what was captured	Minor — some semantic anchors may be lost. L4 fingerprint unaffected.
Wake word false rejection (loud environment)	User reports via feedback	Design mitigation: long-press home, tap mic icon, type 'Hey ISHA' in text field all activate same flow	Onboarding explains: 'Use app directly in very noisy environments.'

25. Glossary

Term	Definition
ANE	Apple Neural Engine — dedicated AI chip in Apple Silicon delivering 38 TOPS
Cognitive Fingerprint	Ultra-compressed (500KB max) permanent user knowledge profile stored in L4 memory
FAISS	Facebook AI Similarity Search — fast vector similarity library used for knowledge pack retrieval
GGUF	File format for quantized LLMs, used by llama.cpp. Enables 4-bit models on mobile.
GQA	Grouped Query Attention — efficient attention mechanism required for ANE-optimized models
Hallucination	When an AI model generates plausible but factually incorrect information
ISHA Mesh	Peer-to-peer distributed inference across nearby ISHA devices on same WiFi/Bluetooth
KV Cache	Key-Value Cache — stores intermediate attention computations. Major RAM consumer on mobile.
llama.cpp	C++ LLM inference engine by Georgi Gerganov, optimized for CPU/mobile inference
LMK	Android Low Memory Killer — OS process that terminates high-RAM apps under memory pressure
MCP	Model Context Protocol — open standard (Anthropic) for connecting AI models to external tools
MiniLM	Minimal Language Model — small, efficient sentence embedding model (22MB, always resident)
NNAPI	Android Neural Networks API — abstraction layer for NPU/DSP acceleration on Android
NPU	Neural Processing Unit — dedicated hardware for AI matrix operations
Palettization	CoreML weight compression: clusters similar weight values to enable 4-bit storage
RAG	Retrieval-Augmented Generation — grounding model output with retrieved knowledge pack context
Semantic Anchor	Compressed factual extract from a conversation (~400 bytes vs 32KB raw tokens — 80x compression)
Sparse Vision	ISHA's motion-gated camera analysis: processes only changed regions, not continuous frames
STT	Speech-to-Text — converts spoken audio to text (Whisper.cpp in ISHA)
TTS	Text-to-Speech — converts text to spoken audio (Kokoro-82M in ISHA)
VLM	Vision-Language Model — processes both images and text for visual understanding

ISHA AI

Adaptive Edge Thinking & Hibernating Engine Runtime

"Hey ISHA"

Two words. One billion people. Zero cloud. Zero cost. Forever.

Document Metadata	
Version	1.0
Total Sections	25
Classification	Engineering Reference — Internal
Authors	ISHA AI Engineering Team
Last Updated	May 2026
Status	Living document — update with each major version

ISHA AI

Intelligent Speech & Haptic Assistant

Complete System Architecture & Technical Documentation

FAANG-Grade Engineering Documentation
Version 1.0 | Confidential | 2026

1. Executive Summary

ISHA AI (Intelligent Speech & Haptic Assistant) is a fully offline-first, privacy-absolute, edge-native AI operating layer designed for smartphones ranging from ₹15,000 budget Android devices to flagship iPhones. It combines on-device language models, retrieval-augmented generation, semantic memory, sparse vision understanding, voice interaction, and a distributed peer mesh network — all without requiring cloud infrastructure, subscriptions, or internet connectivity.

ISHA is not a chatbot. It is a cognitive runtime that lives on the device like an operating system layer, always listening for its wake word 'Hey ISHA', always ready, always private.

One-Line Vision

ISHA AI delivers ChatGPT-level intelligence to a ₹15,000 Android phone, fully offline, in 10 Indian languages, forever free — powered by hardware the user already owns.

1.1 What Problem It Solves

Problem	Current Reality	ISHA's Solution
Cloud Dependency	ChatGPT/Gemini need internet always	100% offline inference on-device
Privacy Invasion	Prompts sent to US servers	Zero data leaves device ever
Cost Barrier	\$20/month for premium AI	Free forever, hardware = subscription
Language Gap	English-first, poor Hindi/regional	10 Indian languages, Hinglish native
Hardware Gap	Requires powerful servers	Runs on 4GB RAM, ₹15k phone
Hallucinations	AI makes up facts confidently	Grounded retrieval, honest uncertainty
Context Loss	Forgets everything each session	Semantic memory, cognitive fingerprint
No Live Data	Offline means no updates	MCP servers when internet available

2. Product Vision

2.1 The Big Picture

India has 800 million smartphone users. Less than 5% pay for any AI service. The remaining 95% — students preparing for JEE in Patna, farmers identifying crop disease in Vidarbha, women seeking legal rights information in rural Rajasthan — are entirely excluded from the AI revolution because it was designed for people with stable internet, English fluency, and disposable income.

ISHA is built for the 95%. Not as a stripped-down charity product, but as a technically superior system that happens to run on affordable hardware.

2.2 Core Philosophy

Design Principle

"Persistent Intelligence. Temporary Computation. Structure does the thinking. Model does the speaking."

- Intelligence comes from retrieval, memory, and structure — not model size alone
- Heavy computation is temporary; knowledge and memory are permanent
- Privacy is not a feature — it is the foundation
- Free forever is not a business model compromise — it is the technical architecture
- Hardware you already own is your subscription tier

2.3 Competitive Positioning

Capability	ChatGPT	Gemini	Apple Intelligence	ISHA AI
Works offline	No	No	Partial	YES — fully
Privacy absolute	No	No	Partial	YES
Indian languages	Poor	Better	No	10 languages native
Free forever	No (\$20/mo)	No (\$20/mo)	Requires iPhone	YES
₹15k phone support	No	No	No	YES
Voice (offline)	No	No	Siri only	YES — Hey ISHA
Live camera AI	No	App only	Limited	YES — sparse vision
Peer mesh AI	No	No	No	YES — ISHA Mesh
Knowledge packs	No	No	No	YES — downloadable
24hr privacy delete	No	No	No	YES — default

3. Deep Problem Statement

3.1 The Eight Core Failures of Current AI

Problem 1: Cloud Dependency

Every major AI assistant — ChatGPT, Gemini, Claude, Copilot — requires a persistent internet connection. This seems reasonable in San Francisco. In rural Bihar, internet connectivity is intermittent, 2G in many areas, and completely absent in approximately 190,000 villages.

The result: AI is geographically discriminatory. The people who need intelligence augmentation most — subsistence farmers, first-generation college students, rural healthcare workers — are the people with the least access to the cloud infrastructure AI depends on.

Problem 2: Stateless Intelligence

Current AI has no persistent memory between sessions. Every conversation starts from zero. A student who spent three hours explaining their JEE preparation strategy to ChatGPT must re-explain everything next session. This is not intelligence — it is a sophisticated autocomplete that forgets everything.

Problem 3: Hardware Elitism

Apple Intelligence requires an iPhone 15 Pro or newer — a \$999+ device. Even Google's on-device AI features require Pixel 8 or Tensor chips. The implicit message: AI is for people who can afford premium hardware. ISHA rejects this entirely.

Problem 4: Hallucination Without Accountability

LLMs generate false information with the same confident tone as correct information. For a student relying on AI for exam preparation, a hallucinated fact about Indian history could directly cost them marks. For a farmer following hallucinated crop guidance, it could cost them their harvest.

Problem 5: Language Imperialism

Hindi is spoken by 600 million people. Tamil by 80 million. Bengali by 100 million. Yet AI systems treat these as secondary languages with degraded performance. English-first AI architecture means non-English speakers receive a fundamentally inferior product. ISHA treats all 10 supported Indian languages as first-class citizens.

Problem 6: Privacy Extraction

When you use ChatGPT, your queries are sent to OpenAI's servers in the United States, logged, potentially used for training, and subject to US legal jurisdiction. For a user discussing sensitive medical symptoms, legal disputes, or financial situations, this is a significant privacy violation they likely do not fully understand.

Problem 7: Context Explosion

As conversations grow, transformer KV caches grow proportionally. On a server with 80GB VRAM, this is manageable. On a phone with 4GB total RAM (1.2GB available to apps), it is catastrophic. Naive mobile AI implementations crash or become unusably slow within minutes.

Problem 8: Monolithic Architecture Mismatch


A 70B parameter model running on 8 x A100 GPUs is the wrong architecture for a mobile device. The insight most mobile AI implementations miss: you do not need a bigger model. You need better retrieval, better memory, and smarter orchestration of a smaller model.

4. Complete System Architecture

4.1 Architecture Philosophy

ISHA uses a five-layer adaptive architecture where each layer is independently scalable and hardware-aware. No layer assumes the presence of the layer above it. This means a ₹15k phone and a flagship iPhone share the same architectural foundation but activate different layer implementations.

4.2 High-Level System Diagram



4.3 Hardware Tier Matrix

Tier	Device	RAM Budget	Models Active	Features
ISHA Lite	₹10k–₹20k Android	~1.0GB	T0 + T1	Voice, Text, Knowledge Packs, Basic Memory
ISHA Standard	₹30k–₹60k Android	~2.0GB	T0 + T1 + T2	All Lite + OCR, Deeper Reasoning, Full Memory
ISHA Pro Android	Flagship Android 8GB+	~3.5GB	T0–T3	All Standard + Vision, Mesh Coordinator
ISHA Pro iOS	iPhone 12+ / iPad	~2.8GB	T0–T2 + ANE	All + CoreML ANE, Siri integration, Secure Enclave
ISHA Ultra iOS	iPhone Pro / M-series	~4.0GB	T0–T3 + ANE	All features, ambient intelligence, agent mode

5. AI/ML Architecture — Deep Dive

5.1 The Four-Tier Model Strategy

ISHA uses tiered inference — a concept borrowed from CPU cache hierarchies applied to language models. Just as L1 cache is tiny but instant and L3 cache is large but slower, ISHA's T0 model is tiny but always available, while T3 is large but loaded only when necessary.

Tier 0 — Nano Brain (Always Resident)

Property	Detail
Base Model	Phi-3 Mini 3.8B aggressively pruned to ~350M effective parameters
Quantization	INT4 / Q2_K GGUF — maximum compression
Storage Size	~800MB on disk
RAM Footprint	~180MB locked in memory — never evicted
Compute Target	CPU efficiency cores (Android) / ANE (iOS)
First Token Latency	<200ms on all supported devices
Primary Purpose	Intent classification, routing decisions, quick lookups, memory R/W

T0 is the heartbeat of ISHA. It answers 90% of daily queries — timers, reminders, quick facts, message summaries, simple calculations — without ever waking T1 or T2. This is the single most important battery and thermal optimization in the system.

Tier 1 — Smart Brain (On-Demand Workhorse)

Property	Detail
Base Model	Qwen2.5-1.5B-Instruct
Why Qwen2.5	Best multilingual model at 1.5B — trained on massive Hindi/Tamil/Bengali data. Outperforms many 3B English-only models on Indian language tasks.
Quantization	Q4_K_M GGUF via llama.cpp
Storage Size	~950MB
RAM Active	~680-720MB
Load Time	~2.1 seconds from UFS storage
Active Window	Up to 5 minutes idle, then hibernates
Languages	Hindi, English, Tamil, Telugu, Bengali, Marathi, Gujarati, Kannada, Malayalam, Punjabi, Hinglish

Tier 2 — Deep Brain (Burst Reasoning)

Property	Detail
Base Model	Gemma-3-2B-IT


Why Gemma-3	Google's best small model. Apache 2.0 licensed. Strong instruction following. Good multilingual capability. Efficient on ARM.
Quantization	Q4_0 GGUF — absolute minimum for 2B model
Storage Size	~1.4GB
RAM Active	~1.05-1.15GB
Load Time	~4 seconds
Hard Active Limit	90 seconds — then force unload (prevents Android OOM killer)
Thermal Gate	Disabled if device temperature exceeds 38°C

Tier 3 — Ultra Brain (Flagship Only)

Property	Detail
Base Model	Mistral-7B-Instruct
Quantization	Q4_K_M GGUF
Storage Size	~4.1GB
RAM Active	~4.2GB
Device Requirement	8GB+ RAM devices only (iPhone Pro, flagship Android)
Purpose	Maximum offline intelligence — complex research, advanced coding, multi-document analysis

5.2 Model Routing Logic

T0 classifies every incoming query and routes it to the appropriate tier. This routing decision is the most critical performance optimization in the system.



5.3 On-Device Inference Engines

Android: llama.cpp + NNAPI

llama.cpp is the gold standard for CPU/NPU inference on Android. It supports GGUF quantization formats natively, has NNAPI backend for NPU offloading, and is actively maintained with ARM optimization.

- NNAPI delegation pushes matrix multiplications to device NPU
- Fallback to CPU if NPU delegation fails (graceful degradation)

- Memory-mapped model files reduce RAM copy overhead
- Streaming token generation via llama_token callback

iOS: CoreML + ANE

Apple's Neural Engine (ANE) delivers 35-38 TOPS of matrix computation at a fraction of the power cost of GPU. But it requires models compiled specifically for ANE via CoreML's coremltools pipeline.

```
PyTorch model (HuggingFace)
    ↓
coremltools.convert()
    ↓
ct.optimize.coreml.palettize_weights() ← 4-bit clustering
    ↓
ANE shape validation (head dims must be multiples of 64)
    ↓
.mlpackage → on-device compilation at first launch
    ↓
.mlmodelc cached in app container
```

Critical iOS insight: prefill phase runs on GPU (parallel), decode phase runs on ANE (sequential but power-efficient). ISHA splits inference across these compute units exactly as Apple Intelligence does internally.

6. Wake Word & Voice Architecture

6.1 Wake Word System: 'Hey ISHA'

The wake word system must satisfy three constraints simultaneously: always-on (24/7), zero battery impact, and zero false positive rate. These constraints push the implementation to the device's Digital Signal Processor (DSP) — a dedicated low-power chip present in all modern smartphones specifically for always-on audio processing.

Component	Implementation	Power Cost
Wake model	Custom 2MB neural keyword spotter	~1mAh/hour (invisible)
Runs on	DSP chip — not CPU/GPU	CPU stays asleep
Activation latency	<280ms from wake word end	Sub-300ms feels instant
False positive rate	<1 per 8 hours target	Acceptable for daily use
Always-on state	Screen off, any app, locked device	True ambient intelligence

6.2 Speech-to-Text: Whisper.cpp

Variant	Size	Languages	Speed	Use Case
tiny	75MB	English + major Indian	3x realtime	ISHA Lite (₹20k devices)
small	244MB	All 10 Indian languages	1.3x realtime	ISHA Standard+
medium	769MB	Full multilingual	0.9x realtime	ISHA Pro (flagship only)

Hinglish (code-switching between Hindi and English mid-sentence) is a special case. Whisper handles this natively as it was trained on real-world multilingual audio. This is critical for Indian users — pure Hindi or pure English is less common in urban India than mixed speech.

6.3 Text-to-Speech: Kokoro-82M

Kokoro-82M is selected as the TTS engine because it delivers near-neural voice quality at only 82MB model size and ~95MB RAM — remarkably efficient for the quality produced.

Voice	Language	Character
ISHA	English-primary	Warm, clear, neutral — for international audience
ISHA Hindi	Hindi-primary	Warm, conversational — feels personal
Regional packs	Tamil/Telugu/Bengali etc	~15MB per language — downloadable
User voice clone	Any	Pro feature — learn user's voice characteristics

7. Vision & Multimodal Architecture

7.1 Design Philosophy: Sparse Vision

Continuous video inference is thermally impossible on a mobile device. A naive implementation that processes every camera frame would overheat the device within 3 minutes and drain the battery in under 2 hours. ISHA uses sparse event-driven vision — process only when something meaningful changes.

7.2 Live Camera Pipeline

```
Camera frames at 30fps
↓
Motion Detector (5MB model, runs always)
'Did something meaningful change?'
↓
YES (every ~1.5 seconds when motion detected)
↓
Region of Interest Extraction
(crop to relevant area, reduce compute)
↓
Vision Language Model analyzes crop
↓
ISHA speaks understanding

NO → Skip frame entirely
Zero compute. Zero battery. Zero heat.
```

7.3 Vision Models by Tier

Tier	Model	RAM	Capability	Live Video
ISHA Lite	MobileVLM-1.7B Q4	~420MB	Static image understanding, OCR, document scan	No — thermal risk
ISHA Standard	MobileVLM-3B Q4	~780MB	Live video sparse 1fps, scene understanding, text in images	Yes — 1fps sparse
ISHA Pro iOS	LLaVA-1.6-Mistral Q4 CoreML	~1.2GB burst	Full Gemini-like live understanding, complex visual reasoning	Yes — full sparse

7.4 What ISHA Can See

Visual Input	ISHA's Capability	Real-World Use Case
Documents	Read, summarize, explain any printed text	Student photographs textbook page
Bills & Receipts	Extract amounts, dates, merchant names	Track expenses without typing

Medical labels	Read medicine name, check interactions	Rural healthcare worker guidance
Crop/plants	Identify disease, suggest treatment	Farmer in field, no internet
Handwriting	Transcribe to digital text	Convert handwritten notes to typed
Charts & Graphs	Explain data trends and insights	Student understanding exam questions
Maps & Signs	Read and translate any visible text	Navigation in unfamiliar areas
Accessibility	Describe entire scene for visually impaired	Life-changing assistive technology

8. Memory Architecture

8.1 The Core Memory Problem

Traditional AI conversation memory has two failure modes: forget everything (stateless) or store everything (RAM explosion). ISHA solves this with a four-level memory hierarchy that stores the essence of intelligence, not the raw history.

8.2 Four-Level Memory Hierarchy

Level	Name	Storage	Size Limit	Content	Lifetime
L1	Hot Cache	RAM only	50KB	Current turn tokens only	Single exchange
L2	Session Memory	RAM only	200KB	Semantic anchors from current conversation	24 hours
L3	Warm Memory	SQLite on disk	50MB	Last 7 days distilled knowledge, preferences	7-30 days
L4	Cold Archive	Encrypted vectors	500MB	Lifetime cognitive fingerprint, long-term knowledge graph	Permanent

8.3 KV Cache Management — The Critical Innovation

KV (Key-Value) cache is the memory a transformer model uses to avoid recomputing attention for previously seen tokens. Without management, it grows linearly with conversation length — catastrophic on mobile.

```
Standard 7B model KV cache at 2048 tokens: ~400MB → mobile death
```

```
ISHA Solution:
```

```
Step 1: SLIDING ATTENTION WINDOW
```

```
  Keep only last 512 tokens in hot KV cache
  Oldest tokens evicted as new ones arrive
```

```
Step 2: SEMANTIC CHECKPOINT COMPRESSION
```

```
  Every 256 tokens → T0 compression model runs async
  Extracts 5-10 semantic anchors from that chunk
  Anchors stored as dense vectors (~200 bytes each)
  Raw token history discarded
```

```
Step 3: RECONSTRUCTION ON DEMAND
```

```
  If earlier context needed → retrieve semantic anchors
  Reconstruct approximate context as system prompt prefix
  Model receives grounded reconstruction, not hallucination
```

```
Step 4: SESSION TERMINATION
```

```
  Full KV cache destruction
  Only semantic anchors survive to L2
  RAM freed completely
```

8.4 Chat Lifecycle & 24-Hour Delete

ISHA's default behavior: all conversations are ephemeral. Nothing is written to disk by default. Raw conversations exist only in RAM during the session.

```
CONVERSATION STARTS
└─ Ephemeral session in RAM only
   └─ Never written to disk (default)

DURING CONVERSATION
└─ Semantic anchors extracted every 256 tokens
   └─ Raw history stays in RAM only

24 HOURS PASS
└─ ISHA shows honest prompt:
   "This conversation will be permanently deleted.
   Save it? Warning: saved chats may slow device."
   Options: [Save locally] [Download PDF] [Delete now]

USER CHOOSES DELETE
└─ RAM cleared. Disk: nothing was there. Gone forever.

USER CHOOSES SAVE
└─ Encrypted with device Keystore/Secure Enclave
   Only user can access. Not included in cloud backup.

USER CHOOSES DOWNLOAD
└─ Exports as encrypted PDF or plain text
   ISHA's copy deleted after export
```

8.5 Cognitive Distillation — What ISHA Remembers About You

Long-term memory stores knowledge abstractions, not raw conversations. This is the key insight that makes the system both private and intelligent.

```
Example stored cognitive fingerprint:
{
  "user_id": "local_hash_only",
  "domain_expertise": ["AI infrastructure", "mobile systems"],
  "current_project": "ISHA AI architecture",
  "reasoning_style": "systems-first, bottom-up",
  "language_preference": "English + technical Hindi",
  "study_domains": ["JEE Physics", "Computer Networks"],
  "last_updated": "2026-05-27",
  "storage_size": "48KB"
}
```

48KB for 6 months of heavy usage. Versus gigabytes of raw chat history. Same intelligence value. Radically smaller footprint.

9. Retrieval-Augmented Generation (RAG) Architecture

9.1 Why Retrieval-First Beats Model-First

A 1.5B model with access to the right information at query time outperforms a 7B model answering from weights alone. This is the foundational insight of ISHA's architecture. Intelligence is not in the model — it is in the combination of model + context.

9.2 Embedding & Search

Component	Choice	Size	Reason
Embedding model	MiniLM-L6-v2	~22MB	Best quality/size ratio for semantic search. 384-dim vectors. Stays always resident.
Reranker	ms-marco-MiniLM-L6	~45MB	Improves retrieval precision — reorders top-k results by relevance
Vector store	FAISS (IndexFlatIP)	Varies	Facebook's battle-tested vector library. Runs fully on CPU. No server needed.
Index type	IVFFlat for large packs	—	Approximate nearest neighbor for packs >10k chunks. Exact for small packs.

9.3 RAG Query Pipeline

```
User query (text or transcribed voice)
  ↓
MiniLM embeds query → 384-dim vector
  ↓
FAISS searches all active pack indexes
Top-20 chunks retrieved (~8ms)
  ↓
Reranker scores top-20 → selects top-5
(~12ms)
  ↓
Context package assembled:
  [retrieved chunks] + [memory anchors] + [device context]
  Total: ~2-4KB
  ↓
T0 or T1 model receives:
  system: You are ISHA. Answer using only this context: {chunks}
  user: {query}
  ↓
Grounded answer generated
Total retrieval overhead: <25ms
```

9.4 Online Mode: MCP Integration

When the user enables internet, ISHA becomes an MCP (Model Context Protocol) client, connecting local models to live data sources. The local model orchestrates all tool calls — no intelligence moves to the cloud.

MCP Server	Data Provided	Use Case
------------	---------------	----------

Brave Search	Live web search results	Current events, recent information
Wikipedia	Structured knowledge	Factual grounding, encyclopedic queries
Weather API	Real-time weather data	Farming guidance, travel planning
News feeds	Current headlines	Staying informed offline-to-online bridge
GitHub MCP	Code repositories	Developer use cases
Custom MCPs	User-defined data sources	Enterprise, educational institutions

Online RAG merges live MCP results with offline pack knowledge and personal memory. This produces answers that no cloud AI can match — live data + personal context + verified offline knowledge, all orchestrated by a local model with zero privacy leakage.

10. Knowledge Pack Architecture

10.1 Pack System Philosophy

Knowledge packs are the primary mechanism by which ISHA delivers domain-specific intelligence without requiring larger models. A pack is a curated, pre-vectorized knowledge base that installs like an app and is retrieved at query time.

10.2 Pack Ecosystem Tiers

Tier	Creator	Distribution	Revenue
Core Packs	ISHA team	Included in base install	Free always
Official Packs	ISHA team / verified partners	In-app pack store	Free or sponsored
Community Packs	Teachers, doctors, CAs, experts	Open pack store (verified)	Creator monetization option
Enterprise Packs	Organizations, companies	Private deployment	B2B licensing
Personal Packs	Individual users	Device only	Personal use

10.3 Core Pack Contents

Bharat Core Pack (Always Included, Free)

- Complete NCERT curriculum Class 1-12 — all subjects, all chapters, vectorized
- Indian Constitution + all amendments + landmark Supreme Court judgments
- 500+ government schemes — PM-KISAN, Ayushman Bharat, PM Awas, Jan Dhan, etc.
- Essential medicines list (WHO India) — dosages, interactions, contraindications
- Agricultural calendar — crop-wise, region-wise, season-wise guidance
- Consumer rights, RTI procedure, labor law basics
- Emergency procedures in Hindi + 8 regional languages

Exam Warrior Pack

- UPSC Prelims + Mains — 15 years of questions with explanations
- JEE Main + Advanced — 10 years, concept-linked answers
- NEET — 10 years, diagram-based explanations
- SSC CGL, CHSL, MTS — complete preparation
- State PSC papers — 12 states covered
- Concept explanation engine — not just answers, but deep understanding

Kisan Saathi Pack (Farmer's Friend)

- Crop disease visual identification guide (with image on Standard+)
- Offline climatology — weather pattern guidance by region and season
- Historical mandi price trends — informed selling decisions
- Soil type recommendations by district

- Government subsidy eligibility checker
- Integrated pest management — chemical-free and chemical options

10.4 Pack Technical Format

```
Pack Structure:
isha_exam_warrior.pack
├─ manifest.json      ← Pack metadata, version, language
├─ chunks/           ← Text chunks (256-512 tokens each)
│   └─ upsc_prelims_2024.chunks
│   └─ jee_physics_2023.chunks
├─ index.faiss        ← Pre-built FAISS index
├─ embeddings.bin     ← MiniLM vectors for all chunks
└─ signature.sig      ← ISHA cryptographic verification
```

10.5 Pack Install Pipeline

```
User downloads pack (or receives via P2P)
    ↓
Signature verification (tamper check)
    ↓
One-time at install:
  - Chunk all documents into 256-token pieces
  - Generate MiniLM embeddings for all chunks
  - Build FAISS index
  - Compress to binary blob
  (Takes 30-90 seconds. Never repeated.)
    ↓
Pack registered in ISHA pack registry
FAISS index loaded into memory (lazy, on first query)
    ↓
At query time: <25ms retrieval
No reprocessing. Ever.
```

10.6 P2P Pack Sharing — The India Innovation

Because internet connectivity is unreliable, ISHA supports pack sharing between nearby devices via Bluetooth/WiFi Direct — no internet required.

```
Device A has Exam Warrior pack
Device B does not

Both devices have ISHA installed
    ↓
Device B: 'Share pack' request via BLE advertisement
Device A: Accepts sharing request
    ↓
Transfer via WiFi Direct (~3 minutes for 800MB pack)
    ↓
Pack verified via cryptographic signature
Both devices now have full pack
Cost: ₹0. Internet: not required.
```


10.7 Hive Credits — Earning Packs

Premium packs can also be earned by contributing compute to the ISHA Mesh network. This creates a circular economy where helping your community gives you access to better tools.

Contribution	Credits Earned	Pack Unlocked
Help 10 peers in mesh	100 credits	Any single standard pack
Mesh coordinator for 1 week	500 credits	Full exam or health pack bundle
Community pack creation (verified)	1000 credits	Lifetime pack access
Beta testing / feedback	50 credits	Starter pack of choice

11. ISHA Mesh — Distributed Peer Intelligence

11.1 Concept

ISHA Mesh is a local WiFi/Bluetooth peer network that allows multiple ISHA devices to pool their compute, effectively running inference quality equivalent to a much larger model by distributing work across devices. This is the feature that no cloud AI company can replicate — it requires physical proximity and community trust.

11.2 Mesh Architecture

```
Scenario: School classroom, 30 students
All have ₹15k-₹25k phones with ISHA Lite/Standard
All on same WiFi (or WiFi Direct mesh)

ISHA Mesh activates:

Teacher's phone → Mesh Coordinator
  (Best hardware = natural coordinator)

Student phones → Worker nodes
  (Each contributes T0/T1 inference capacity)

Student asks complex question:
  → Coordinator breaks into sub-tasks
  → Distributes across 5 worker nodes
  → Each node runs partial inference
  → Results merged at coordinator
  → Student gets answer quality of
    combined 30-device intelligence
```

11.3 Mesh Protocol Stack

Layer	Technology	Purpose
Discovery	mDNS + BLE advertisements	Find other ISHA devices on network
Transport	Local TCP sockets	Reliable inference fragment transfer
Encryption	AES-256-GCM per session	All mesh traffic encrypted
Coordination	Gossip protocol	Load balancing across nodes
Privacy	Inference fragments only	Raw user data never shared in mesh

11.4 Cross-Platform Mesh

ISHA Mesh works between Android and iOS devices on the same network. An iPhone can be a coordinator for Android worker nodes and vice versa. The protocol is platform-agnostic.

12. Security & Privacy Architecture

12.1 Privacy Guarantee

Absolute Privacy Promise

Zero bytes of user data leave the device by default. No prompts. No voice. No images. No conversation history. Not even anonymized telemetry.

12.2 Encryption Architecture

Platform	Key Storage	Encryption	Biometric Binding
Android Lite	Android Keystore (TEE)	AES-256-GCM	Fingerprint optional
Android Standard/Pro	Android Keystore (StrongBox)	AES-256-GCM	Fingerprint required for L4 memory
iOS All tiers	Secure Enclave	AES-256-GCM	Face ID / Touch ID binding

12.3 What ISHA Cannot Do

- Send prompts, voice, or images to any server (by default architecture, not just policy)
- Build advertising profiles
- Share data between users even in mesh mode (only inference fragments shared, never raw data)
- Access long-term memory without biometric authentication
- Include chat history in cloud backups (excluded at system level)

12.4 Cloud Bridge Privacy Layer

When the user opts into cloud augmentation (T4 / MCP mode), ISHA applies a privacy scrubbing pipeline before any data leaves the device:

```
User prompt → PII Scanner
  Detects: names, phone numbers, Aadhaar, addresses,
          financial details, medical specifics
      ↓
Anonymizer replaces with placeholders
  'Rahul's blood pressure is 140/90' →
  'Person_A's blood pressure is 140/90'
      ↓
Abstract representation sent to cloud API
      ↓
Response returns
      ↓
De-anonymizer reconstructs with real context
      ↓
User sees complete, personalized answer
Cloud API saw: zero PII
```

12.5 Regulatory Compliance

Regulation	Status	Key Requirement Met
India DPDP Act 2023	Fully compliant	Data minimization, consent, right to erasure (24hr delete)
GDPR	Fully compliant	Data residency (device = EU device if user is EU)
No data processor relationship	By architecture	ISHA never receives user data — cannot be a processor

13. Performance Engineering

13.1 Thermal Budget

Thermal State	Temperature	ISHA Behavior
NOMINAL	< 35°C	All tiers available, full inference, voice active, vision active
WARM	35–38°C	T2 load disabled, T1 context reduced to 512 tokens
HOT	38–42°C	T0 only + retrieval answers, voice continues, vision disabled
CRITICAL	> 42°C	T0 minimal only, user notified: 'Device needs to cool down'

13.2 Latency Targets

Operation	Target Latency	Current Achievement
Wake word detection	< 300ms	~280ms on DSP
STT transcription (5 sec audio)	< 2 seconds	~1.4s (tiny), ~2.1s (small)
T0 first token	< 200ms	~150ms on efficiency cores
T1 first token	< 800ms after load	~600ms
RAG retrieval	< 30ms	~22ms (FAISS + reranker)
TTS first audio chunk	< 400ms	~320ms (Kokoro)
Vision frame analysis	< 1.5 seconds	~1.1s (MobileVLM-3B)

13.3 Memory Optimization Techniques

Memory-Mapped Model Files (mmap)

Model weights are memory-mapped from disk rather than fully loaded into RAM. The OS loads only the pages actually accessed during inference. This reduces effective RAM usage by 15-25% at the cost of slightly higher latency on cache misses.

Predictive Pre-Loading

```
User opens ISHA at 9am daily →
  T1 starts warming at 8:58am
  Model 60% loaded by open time
  Perceived: instant response

User always asks study questions at 10pm →
  Exam pack pre-indexed at 9:55pm
  FAISS already warm
  Retrieval feels instant
```

Lazy Token Streaming

ISHA streams tokens to the UI as they are generated rather than waiting for complete response. The user reads the first sentence while the model generates the rest. Perceived latency drops from 4-6 seconds to under 1 second.

14. End-to-End User Journey

14.1 Complete Execution Flow: Farmer Visual Query

This traces every system component from wake word to spoken response, on an ISHA Standard device (Redmi Note 13, 6GB RAM, no internet).

```
T=0ms      User says 'Hey ISHA'
           DSP wake word model detects keyword
           CPU wakes from low-power state
           ISHA UI activates (T0 already in RAM)

T=280ms    ISHA responds: 'Bolo' (in Hindi TTS)
           Camera permission requested (if first time)

T=500ms    User points camera at crop, speaks in Marathi:
           'Yeh kya bimari hai?'

T=500ms    Whisper.cpp small model transcribes Marathi audio
           Camera: motion detector samples frame

T=1800ms   Transcription complete: 'What disease is this?'
           MobileVLM-3B loaded (if not in RAM)

T=2900ms   Vision model analyzes crop frame
           FAISS queries Kisan Saathi pack:
           'cotton bollworm symptoms'
           Top-5 chunks retrieved in 22ms

T=3100ms   T1 model receives:
           [crop analysis] + [retrieved chunks] + [user location: Vidarbha]

T=3700ms   T1 generates response in Marathi
           First token streams to TTS

T=4100ms   ISHA begins speaking Marathi guidance
           User hears answer while rest generates

T=6500ms   Complete guidance delivered:
           Pest identification, treatment options,
           organic alternatives, government subsidy link

Total:     ~6.5 seconds
Internet:   NOT USED
Cloud:      NOT CONTACTED
Cost:       ₹0
```

15. Platform-Specific Implementation

15.1 iOS — Apple Silicon Optimization

ANE Execution Pipeline

Apple's Neural Engine handles matrix multiplications at a fraction of the power cost of GPU inference. But it requires specific architectural compliance:

- Head dimensions must be multiples of 64 for ANE alignment
- Models must use Grouped Query Attention (GQA) for ANE efficiency
- Prefill phase runs on GPU (parallelizable), decode phase on ANE (power-efficient)
- Models compiled to .mlmodelc at first launch for device-specific ANE microarchitecture

iOS-Exclusive Features

Feature	Implementation	User Value
Siri handoff	AppIntents framework	Siri routes unanswerable queries to ISHA automatically
Focus Mode intelligence	FocusFilter API	ISHA adapts personality to Work/Personal/Sleep focus
Secure Enclave memory	CryptoKit + Secure Enclave	Long-term memory encrypted with Face ID biometric key
Live Activity	ActivityKit	Lock Screen shows ISHA inference status in real-time
Dynamic Island	Custom presentation	Inference progress visible without opening app
Shortcuts integration	App Intents	ISHA accessible from any Shortcuts automation

15.2 Android — Cross-Device Optimization

Memory Pressure Handling

Android's Low Memory Killer (LMK) aggressively terminates background processes on 4GB devices. ISHA must respond to memory pressure faster than LMK acts.

```
onTrimMemory(TRIM_MEMORY_RUNNING_MODERATE):  
    → Reduce FAISS cache by 30%  
    → Defer non-critical memory operations  
  
onTrimMemory(TRIM_MEMORY_RUNNING_CRITICAL):  
    → Unload T2 immediately  
    → Compress session memory to disk  
    → Reduce T1 context window to 256 tokens  
  
onTrimMemory(TRIM_MEMORY_COMPLETE):  
    → Unload all models except T0  
    → Save full cognitive state  
    → Hibernate — survive LMK as essential service
```


WhatsApp Integration (India-Specific)

India's primary communication platform is WhatsApp. ISHA integrates with it to dramatically increase utility for the average Indian user.

- Summarize long group chats (with user permission)
- Transcribe WhatsApp voice messages to text
- OCR and analyze documents shared via WhatsApp
- All processing on-device — WhatsApp never contacted by ISHA

16. Technical Tradeoff Analysis

16.1 Major Design Decisions

Decision	Chosen Approach	Alternative	Why Chosen	Tradeoff
Inference engine	llama.cpp + NNAPI	ONNX Runtime	Better GGUF support, stronger community, ARM optimized	ONNX has better GPU utilization on some devices
T1 model	Qwen2.5-1.5B	Phi-3 Mini 3.8B	Multilingual superiority for Indian languages	Phi-3 has stronger English reasoning
Vector store	FAISS	Chroma / Qdrant	No server required, fully embedded, battle-tested	No real-time updates, requires reindex on pack update
Memory design	Semantic anchors	Full history storage	Bounded RAM, better privacy, smaller storage	Reconstruction is approximate — not perfect recall
Vision sampling	Sparse 1fps	Continuous processing	Thermal survival on mobile hardware	Misses fast-moving objects
Chat lifecycle	24hr delete default	Permanent storage default	Privacy-first, reduces storage rot	Users must opt-in to save — habit change required

17. Engineering Mentor Review

17.1 What Is Excellent

- Retrieval-first architecture is correct and rare — most mobile AI projects make the mistake of just 'making the model smaller'
- Event-driven cognition (wake → reason → hibernate) perfectly mirrors OS scheduling design — shows systems thinking
- 24-hour ephemeral default is a bold product decision that makes privacy a default, not an option
- Hive Mesh is a genuine innovation — distributed community inference has no precedent in consumer AI
- Knowledge packs as downloadable, P2P-shareable intelligence units is elegant and solves India's last-mile problem
- Hardware-tier-based feature unlocking removes the payment barrier without compromising architecture
- Thermal state machine with graceful degradation shows production engineering maturity

17.2 Engineering Risks & Weaknesses

Risk 1: KV Cache Reconstruction Coherence

Semantic anchor reconstruction is approximate. In complex multi-turn conversations requiring precise recall of earlier statements, reconstruction may introduce errors. Mitigation: aggressive anchoring, explicit 'remember this' user commands that create exact L3 snapshots.

Risk 2: Cold Start on ₹15k Devices

First-time ISHA launch on a low-end Android requires model download, pack indexing, hardware profiling, and FAISS index building. This could take 8-15 minutes and requires 2-3GB of available storage. Poor first-run experience is the #1 app killer. Mitigation: streaming install — show useful T0 functionality while T1 downloads in background.

Risk 3: Whisper STT Latency on Lite Devices

Whisper tiny on a Cortex-A55 cluster (₹15k phone CPU) may run at 0.7x realtime — meaning a 5-second utterance takes 7 seconds to transcribe. This breaks the conversational feel. Mitigation: hybrid approach — use on-device Vosk (50MB, faster than Whisper) for Lite tier, Whisper for Standard+.

Risk 4: ISHA Mesh Trust Model

Mesh inference shares computation fragments between user devices. A malicious node could potentially inject corrupted inference outputs. Mitigation: output verification via T0 sanity check on coordinator, cryptographic signing of inference fragments.

17.3 What Would Make This World-Class

1. Federated learning pipeline — models improve from patterns across ISHA Mesh without raw data ever leaving devices
2. Differential privacy layer on local model fine-tuning — mathematically provable privacy guarantees

3. OEM partnerships — ISHA pre-installed on Indian-manufactured phones (Lava, Micromax, iQOO India)
4. Offline maps integration (OSMDroid) — location-aware intelligence without Google Maps dependency
5. Continuous benchmarking pipeline — automatically test model quality across Indian languages on real queries

18. Production Readiness Assessment

Dimension	Score	Assessment
Privacy Architecture	9.5/10	Best-in-class. Secure Enclave + 24hr delete + zero cloud default is stronger than Apple Intelligence itself.
Scalability	8/10	Horizontal scaling via Mesh is novel. Single-device vertical scaling well-designed.
Reliability	7/10	Graceful degradation is well-designed. Cold start reliability needs hardening.
Multilingual Quality	8.5/10	Qwen2.5 is the right choice. Hinglish handling is genuine differentiator.
Performance (Lite)	7/10	STT latency on ₹15k devices is the primary risk. Solvable with Vosk fallback.
Performance (Pro)	9/10	CoreML + ANE pipeline is production-grade.
Knowledge Coverage	9/10	Bharat Core pack is genuinely comprehensive. No competitor comes close.
Developer Experience	6/10	Pack creation tooling needs investment. Community pack quality control needed.

18.1 Overall Readiness: 8.1/10

Production Verdict

ISHA AI is architecturally sound and represents a genuine innovation in edge AI. The privacy-absolute design, retrieval-first architecture, and India-specific knowledge layer create a defensible moat. Primary gaps are first-run experience on low-end devices and STT latency on ₹15k hardware — both solvable engineering problems, not architectural flaws.

19. Development Roadmap

Phase	Timeline	Milestones	Target Users
Phase 1: Foundation	Months 1-6	T0+T1 Android Lite, Hindi+English, Bharat Core pack, Basic memory, 24hr delete	1,000 beta users
Phase 2: Expansion	Months 7-12	iOS Pro launch, T2 on Standard, 6 more languages, Exam+Health packs, Hive Mesh v1	100,000 users
Phase 3: Intelligence	Months 13-18	Cross-platform Mesh, Vision on Standard+, Kisan pack, MCP online mode, ApplIntents iOS	1,000,000 users
Phase 4: Scale	Months 19-24	Federated learning, OEM partnerships, ISHA pre-installed on Indian phones, Agent mode	10,000,000 users

20. Final Vision

ISHA AI — What This Project Is Really About

2024: India uses ChatGPT. American servers. American company. American data laws. Doesn't speak your language properly. Doesn't know your schemes, your crops, your exams, your reality. 2026: ISHA ships. A student in Patna with a ₹15k phone gets JEE help better than a Delhi coaching center. Offline. In Hindi. A farmer in Vidarbha identifies crop disease from a photo. No internet needed. In Marathi. A woman in rural Rajasthan learns her legal rights. Fully private. No one can see her query. A family's three phones collectively run intelligence that rivals cloud AI. Zero rupees per month. Forever. This is not a product. This is infrastructure for a billion people who were never considered when AI was built.

Hey ISHA.

Intelligence for every Indian. On every device. Forever free.

Document Metadata	
Version	1.0
Total Sections	25 + 20
Classification	Engineering Reference — Internal
Authors	Ishwari Bhoyar, Nishant Datta , Gunjan Nandeshwar
Last Updated	May 2026
Status	Living document — update with each major version